

## Motivation

CMOS scaling  $\Rightarrow$   $\uparrow$  device failure  $\Rightarrow$  in-field hardware failure

Need for on-the-fly **detection, diagnosis and recovery/repair**

## Key Observations

Handle only those h/w faults that propagate to s/w  
Optimize common fault-free cases

## Strategy

**Watch for software anomaly (symptoms)**

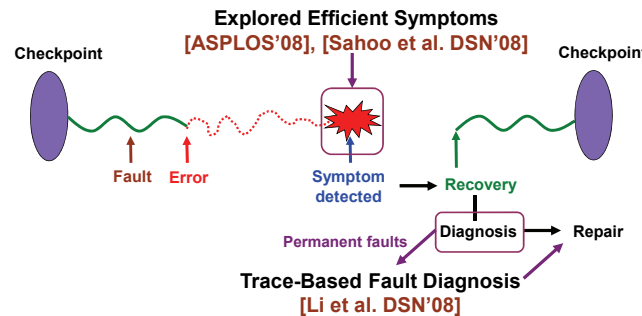
Zero to low overhead "always-on" monitors

**Diagnosis after detection**

May incur high overhead, but *rarely invoked*

**Checkpoint/replay based recovery**

## SWAT: SoftWare Anomaly Treatment



## Ongoing work

### Recovery

**Hybrid hardware/software checkpointing techniques**

Hardware handles most cases

Hybrid technique handles (few) long detection latencies

### SWAT on Multicore Systems

**Understanding fault propagation to fault-free core**

~20% of faults corrupt fault-free core execution

**New symptoms:** No Forward Progress, CPU Panic

## iSWAT [Sahoo et al. DSN'08]

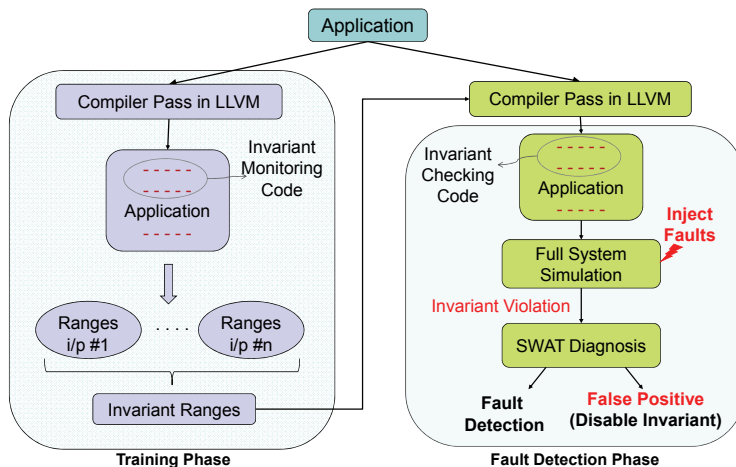
Detection using likely program invariants

Observation:

Undetected faults in SWAT mostly **corrupt data values**

Strategy:

- Check data values with **range-based likely invariants** ( $MIN \leq value \leq MAX$ )
  - Instrument the binary with invariant checking code
- Exploit SWAT diagnosis to **identify false positive** invariants at runtime
  - Disable false positive invariants, limit diagnosis overhead
    - Maximum number of rollback  $\leq$  number of static false positives



## Key Results:

- False Positive rate is less than **5%** with 12 training inputs
- Undetected faults reduced by **28%** and SDCs reduce by **74%**
- Low run-time overhead (5% on x86, 14% on UltraSPARC-III)

## Diagnosis [Li et al. DSN'08]

Symptom-based detection  $\Rightarrow$  distinguish software bugs, transients, permanent faults

Permanent fault needs repair  $\Rightarrow$  disabling core wasteful  $\Rightarrow$  reconfigure faulty **patch units**

### Opportunities

- Exploit **checkpoint/replay** on **multicore** system
- Replay **faulty** and **fault-free** execution, compare
  - $\Rightarrow$  **Synthesized DMR** for diagnosis

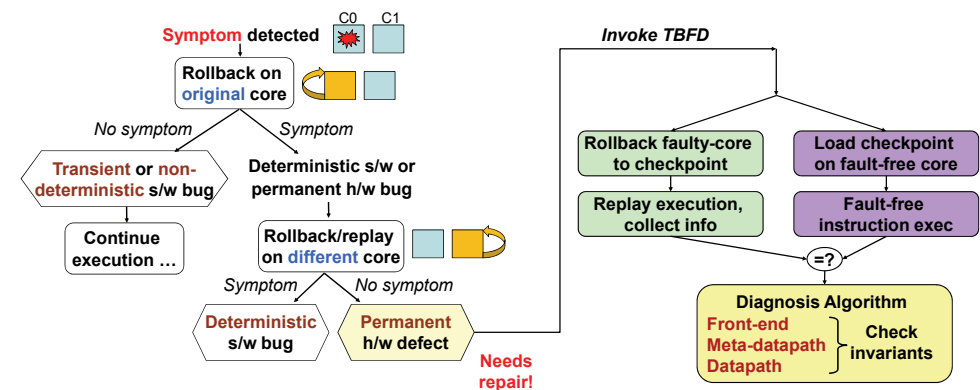
### Trace-Based Fault Diagnosis (TBFD) for patch unit

#### Key Ideas

- Compare instruction trace of faulty vs. fault-free execution
- Divergence  $\Rightarrow$  faulty hardware used  $\Rightarrow$  diagnosis clues

#### Advantages

- Works with other detection techniques
- Supports different repair granularity



## Key Results:

- 98%** of detected faults can be diagnosed
- ~**90%** diagnosed to exact array entry/unit
- Logical-physical register mapping (meta-data path) faults require sophisticated algorithm